

Lecture 5: Multiple Linear Regression

Pilsung Kang

School of Industrial Management Engineering

Korea University

AGENDA

- 01 Multiple Linear Regression
- 02 Evaluating Regression Models
- 03 Variable Selection
- 04 R Exercise

Multiple Linear Regression

- Regression Example: Predict the selling price of Toyota Corolla



Dependent variable
(target)

Independent variables
(attributes, features)

Variable	Description
Price	Offer Price in EUROS
Age_08_04	Age in months as in August 2004
KM	Accumulated Kilometers on odometer
Fuel_Type	Fuel Type (Petrol, Diesel, CNG)
HP	Horse Power
Met_Color	Metallic Color? (Yes=1, No=0)
Automatic	Automatic (Yes=1, No=0)
CC	Cylinder Volume in cubic centimeters
Doors	Number of doors
Quarterly_Tax	Quarterly road tax in EUROS
Weight	Weight in Kilograms


Multiple Linear Regression

- Goal

- ✓ Fit a linear relationship between a quantitative dependent variable Y and a set of predictors X_1, X_2, \dots, X_p .

$$Y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_p x_p + \varepsilon$$

coefficients unexplained



Multiple Linear Regression

- Explanatory vs. Predictive

Explanatory Regression

- Explain relationship between predictors (explanatory variables) and target.
- Familiar use of regression in data analysis.
- Model Goal: Fit the data well and understand the contribution of explanatory variables to the model.
- “goodness-of-fit”: R^2 , residual analysis, p-values.

$$Y = \beta_0 + \beta_1 x_1 + \dots + \beta_p x_p + \varepsilon$$





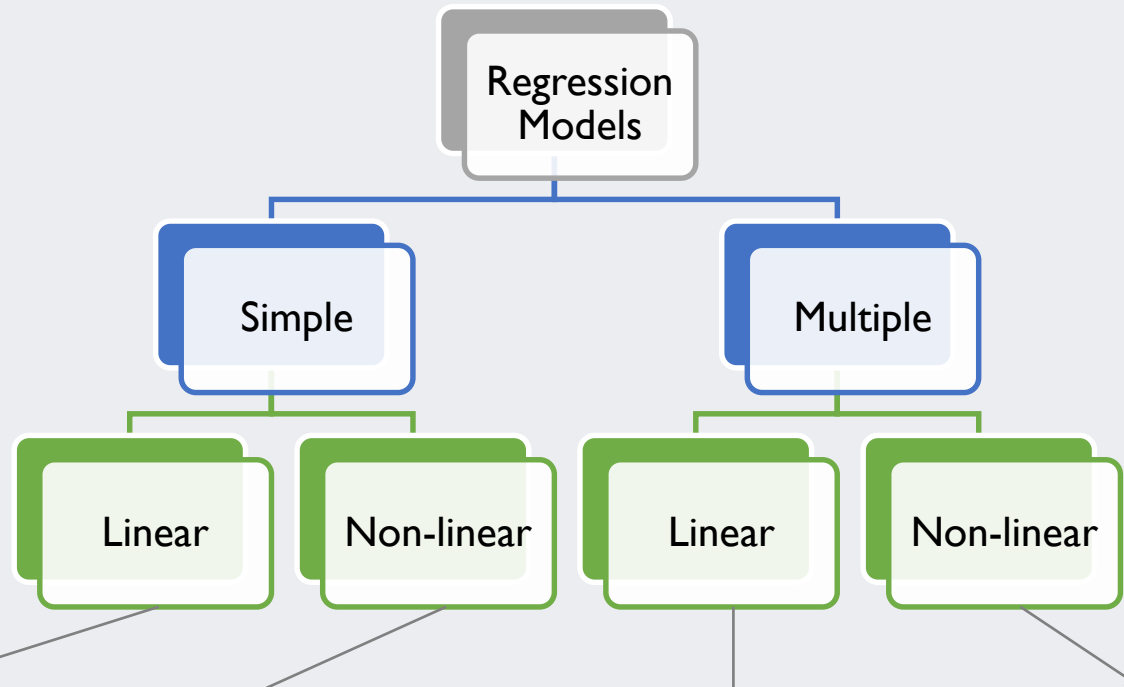
Predictive Regression

- Predict target values in other data where we have predictor values, but not target values.
- Classic data mining context
- Model Goal: Optimize predictive accuracy
- Train model on training data
- Assess performance on validation (hold-out) data
- Explaining role of predictors is not primary purpose (but useful)

$$Y = \beta_0 + \beta_1 x_1 + \dots + \beta_p x_p + \varepsilon$$

Multiple Linear Regression

- Type of Regression




Multiple Linear Regression

- Linear Regression


- ✓ Assume that the relationship between the input variable and the target variable is always **linear**.

$$Y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_p x_p + \varepsilon$$



Multiple Linear Regression

- Which line is optimal?



Multiple Linear Regression

- Estimating the coefficients

- ✓ Ordinary least square (OLS)



- Actual target: $Y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_p x_p + \varepsilon$
 - Predicted target: $\hat{Y} = \hat{\beta}_0 + \hat{\beta}_1 x_1 + \hat{\beta}_2 x_2 + \dots + \hat{\beta}_p x_p$
 - **Goal:** minimize the difference between the actual and predicted target.

$$\begin{aligned} \min \quad \frac{1}{2} \sum_{i=1}^N \varepsilon_i^2 &= \frac{1}{2} (Y_i - \hat{Y}_i)^2 \\ &= \frac{1}{2} \left(Y - \hat{\beta}_0 - \hat{\beta}_1 x_1 - \hat{\beta}_2 x_2 - \dots - \hat{\beta}_p x_p \right)^2 \end{aligned}$$

Multiple Linear Regression

- Estimating the coefficients

✓ Ordinary least square (OLS)



Multiple Linear Regression

- Ordinary least square: Matrix solution

✓ \mathbf{X} : n by p matrix, \mathbf{y} : n by 1 vector, $\boldsymbol{\beta}$: p by 1 vector.

$$\min \quad E(\mathbf{X}) = \frac{1}{2} (\mathbf{y} - \mathbf{X}\boldsymbol{\beta})^T (\mathbf{y} - \mathbf{X}\boldsymbol{\beta})$$

$$\Rightarrow \quad \frac{\partial E(\mathbf{X})}{\partial \boldsymbol{\beta}} = -(\mathbf{y} - \mathbf{X}\boldsymbol{\beta})^T \mathbf{X} = 0$$

$$\Rightarrow \quad -\mathbf{y}^T \mathbf{X} + \boldsymbol{\beta}^T \mathbf{X}^T \mathbf{X} = 0$$

$$\Rightarrow \quad \boldsymbol{\beta}^T = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{y}^T \mathbf{X}$$

$$\Rightarrow \quad \boldsymbol{\beta} = \left((\mathbf{X}^T \mathbf{X})^{-1} \mathbf{y}^T \mathbf{X} \right)^T$$

Multiple Linear Regression

- Ordinary least square

- ✓ Finds the best estimates β when the following conditions are satisfied:

- The noise ε follows a normal distribution.
 - The linear relationship is correct.
 - The cases are independent of each other.
 - The variability in Y values for a given set of predictors is the same regardless of the values of the predictors (homoskedasticity).

Multiple Linear Regression

- Goodness-of-fit: (Adjusted) R^2

$$R^2 = 1 - \frac{\sum_{j=1}^n \hat{\varepsilon}_j^2}{\sum_{j=1}^n (y_j - \bar{y})^2} = \frac{\sum_{j=1}^n (\hat{y}_j - \bar{y})^2}{\sum_{j=1}^n (y_j - \bar{y})^2}$$

- ✓ Gives the proportion of the total variation in the y_i 's explained by the predictor variables
- ✓ R^2 equals 1 if the fitted equation passes through all the data points


Multiple Linear Regression

- Sum-of-Squares Decomposition

$$\sum_{j=1}^n (y_j - \bar{y})^2 = \sum_{j=1}^n (\hat{y}_j - \bar{y})^2 + \sum_{j=1}^n \hat{\varepsilon}_j^2.$$

$\left(\begin{array}{c} \text{total sum of squares} \\ \text{about mean} \end{array} \right) \quad \left(\begin{array}{c} \text{regression} \\ \text{sum of squares} \end{array} \right) \quad \left(\begin{array}{c} \text{residual (error)} \\ \text{sum of squares} \end{array} \right)$

SST **SSR** **SSE**




Multiple Linear Regression

- Coefficient of Determination

- ✓ The proportionate reduction of total variation associated with the use of the predictor variable Z.

$$R^2 = 1 - \frac{SSE}{SST} = \frac{SSR}{SST} \quad 0 \leq R^2 \leq 1$$



Multiple Linear Regression

- Model Fit


- ✓ It is imperative to examine the adequacy of the model before the estimated function becomes a permanent part of the decision making apparatus.

- ✓ For general diagnostic purpose, residuals should be plotted as follows:


1. *Plot the residuals $\hat{\epsilon}_j$ against the predicted values $\hat{y}_j = \hat{\beta}_0 + \hat{\beta}_1 z_{j1} + \cdots + \hat{\beta}_r z_{jr}$. Departures from the assumptions of the model are typically indicated by two types of phenomena:*
2. *Plot the residuals $\hat{\epsilon}_j$ against a predictor variable, such as z_1 , or products of predictor variables, such as z_1^2 or $z_1 z_2$. A systematic pattern in these plots suggests the need for more terms in the model. This situation is illustrated in Figure 7.2(c).*
3. *$Q-Q$ plots and histograms. Do the errors appear to be normally distributed? To answer this question, the residuals $\hat{\epsilon}_j$ or $\hat{\epsilon}_j^*$ can be examined using the techniques discussed in Section 4.6. The $Q-Q$ plots, histograms, and dot diagrams help to detect the presence of unusual observations or severe departures from normality that may require special attention in the analysis. If n is large, minor departures from normality will not greatly affect inferences about β .*

Multiple Linear Regression


- Residual plots




(a)



(b)



(c)




(d)

Multiple Linear Regression

- Model checking


$$y = 2x + \varepsilon, \quad \varepsilon \sim \text{Gamma}(2,1)$$




Regression model



QQ Plot of Residuals



Histogram of Residuals



Multiple Linear Regression: Example

- Example: predict the selling price of Toyota corolla

Y

X

Price	Age_08_04	KM	Fuel_Type	HP	Met_Color	Automatic	cc	Doors	Quarterly_Tax	Weight
13500	23	46986	Diesel	90	1	0	2000	3	210	1165
13750	23	72937	Diesel	90	1	0	2000	3	210	1165
13950	24	41711	Diesel	90	1	0	2000	3	210	1165
14950	26	48000	Diesel	90	0	0	2000	3	210	1165
13750	30	38500	Diesel	90	0	0	2000	3	210	1170
12950	32	61000	Diesel	90	0	0	2000	3	210	1170
16900	27	94612	Diesel	90	1	0	2000	3	210	1245
18600	30	75889	Diesel	90	1	0	2000	3	210	1245
21500	27	19700	Petrol	192	0	0	1800	3	100	1185
12950	23	71138	Diesel	69	0	0	1900	3	185	1105
20950	25	31461	Petrol	192	0	0	1800	3	100	1185
19950	22	43610	Petrol	192	0	0	1800	3	100	1185
19600	25	32189	Petrol	192	0	0	1800	3	100	1185
21500	31	23000	Petrol	192	1	0	1800	3	100	1185
22500	32	34131	Petrol	192	1	0	1800	3	100	1185
22000	28	18739	Petrol	192	0	0	1800	3	100	1185
22750	30	34000	Petrol	192	1	0	1800	3	100	1185
17950	24	21716	Petrol	110	1	0	1600	3	85	1105
16750	24	25563	Petrol	110	0	0	1600	3	19	1065

Multiple Linear Regression: Example

- Data preprocessing

- ✓ Create dummy variables for fuel types

	Fuel_type = Diesel	Fuel_type = Petrol	Fuel_type = CNG
Diesel	1	0	0
Petrol	0	1	0
CNG	0	0	1

- Data partitioning

- ✓ 60% training data / 40% validation data

Id	Model	Price	Age_08_04	Mfg_Month	Mfg_Year	KM	Fuel_Type_Diesel	Fuel_Type_Petrol
1	RRA 2/3-Doors	13500	23	10	2002	46986	1	0
4	RRA 2/3-Doors	14950	26	7	2002	48000	1	0
5	SOL 2/3-Doors	13750	30	3	2002	38500	1	0
6	SOL 2/3-Doors	12950	32	1	2002	61000	1	0
9	VT I 2/3-Doors	21500	27	6	2002	19700	0	1
10	RRA 2/3-Doors	12950	23	10	2002	71138	1	0
12	BNS 2/3-Doors	19950	22	11	2002	43610	0	1
17	ORT 2/3-Doors	22750	30	3	2002	34000	0	1

Multiple Linear Regression: Example

- Fitted linear regression model

Input variables	Coefficient	Std. Error	p-value	SS
Constant term	-3608.418457	1458.620728	0.0137	97276410000
Age_08_04	-123.8319168	3.367589	0	8033339000
KM	-0.017482	0.00175105	0	251574500
Fuel_Type_Diesel	210.9862518	474.9978333	0.6571036	6212673
Fuel_Type_Petrol	2522.066895	463.6594238	0.00000008	4594.9375
HP	20.71352959	4.67398977	0.00001152	330138600
Met_Color	-50.48505402	97.85591125	0.60614568	596053.75
Automatic	178.1519013	212.0528565	0.40124047	19223190
cc	0.01385481	0.09319961	0.88188446	1272449
Doors	20.02487946	51.0899086	0.69526076	39265060
Quarterly_Tax	16.7742424	2.09381151	0	160667200
Weight	15.41666317	1.40446579	0	214696000

β

Significance
Probability

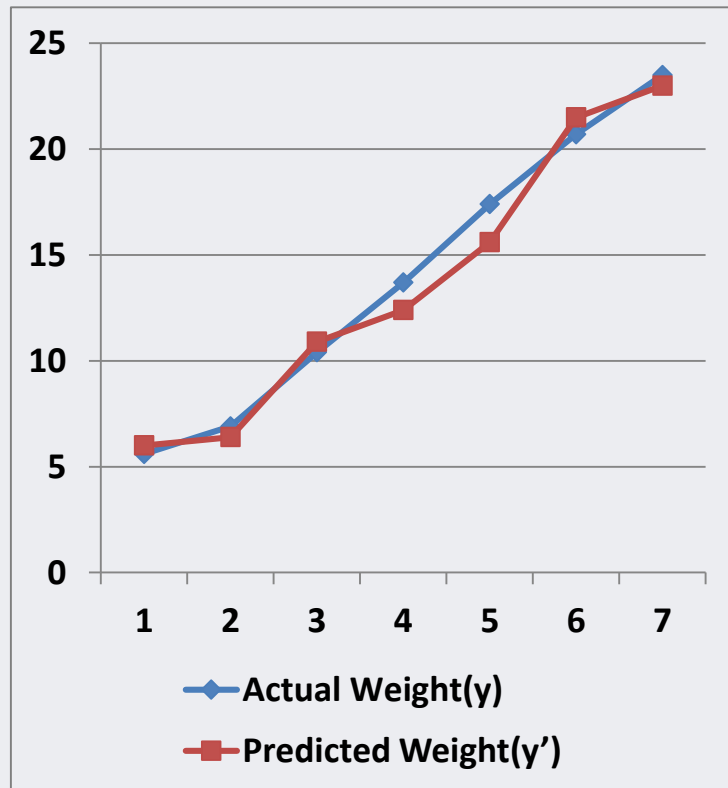
AGENDA

- 01 Multiple Linear Regression
- 02 Evaluating Regression Models
- 03 Variable Selection
- 04 R Exercise

Evaluating Regression Models

- Example: predict a baby's weight (kg) based on his/her age

Age	Actual Weight(y)	Predicted Weight(y')
1	5.6	6.0
2	6.9	6.4
3	10.4	10.9
4	13.7	12.4
5	17.4	15.6
6	20.7	21.5
7	23.5	23.0



Evaluating Regression Models

- Average error

✓ Indicate whether the predictions are on average over- or under-predicted.

Age	Actual Weight(y)	Predicted Weight(y')
1	5.6	6.0
2	6.9	6.4
3	10.4	10.9
4	13.7	12.4
5	17.4	15.6
6	20.7	21.5
7	23.5	23.0

$$\begin{aligned} \text{Average error} &= \frac{1}{n} \sum_{i=1}^n (y - y') \\ &= 0.342 \end{aligned}$$

Evaluating Regression Models

- Mean absolute error (MAE)
 - ✓ Gives the magnitude of the average error

Age	Actual Weight(y)	Predicted Weight(y')
1	5.6	6.0
2	6.9	6.4
3	10.4	10.9
4	13.7	12.4
5	17.4	15.6
6	20.7	21.5
7	23.5	23.0

$$MAE = \frac{1}{n} \sum_{i=1}^n |y - y'|$$
$$= 0.829$$

Evaluating Regression Models

- Mean absolute percentage error (MAPE)

✓ Gives a percentage score of how predictions deviate (on average) from the actual values.

Age	Actual Weight(y)	Predicted Weight(y')
1	5.6	6.0
2	6.9	6.4
3	10.4	10.9
4	13.7	12.4
5	17.4	15.6
6	20.7	21.5
7	23.5	23.0

$$MAPE = 100\% \times \frac{1}{n} \sum_{i=1}^n \frac{|y - y'|}{|y|}$$
$$= 6.43\%$$

Evaluating Regression Models

- (Root) Mean squared error ((R)MSE)
 - ✓ Standard error of estimate
 - ✓ Same units as the variable predicted

Age	Actual Weight(y)	Predicted Weight(y')
1	5.6	6.0
2	6.9	6.4
3	10.4	10.9
4	13.7	12.4
5	17.4	15.6
6	20.7	21.5
7	23.5	23.0

$$MSE = \frac{1}{n} \sum_{i=1}^n (y - y')^2$$
$$= 0.926$$

$$RMSE = \sqrt{\frac{1}{n} \sum_{i=1}^n (y - y')^2}$$
$$= 0.962$$


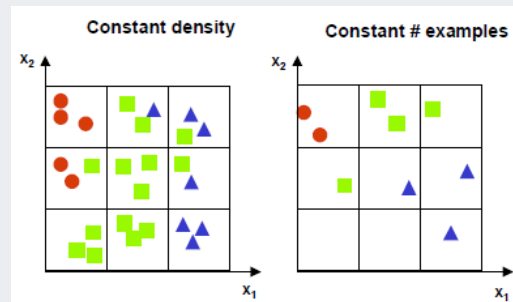


AGENDA

- 01 Multiple Linear Regression
- 02 Evaluating Regression Models
- 03 Variable Selection
- 04 R Exercise

Variable Selection

- Curse of Dimensionality

- ✓ The number of instances increases exponentially to achieve the same explanation ability when the number of variables increases



“If there are various logical ways to explain a certain phenomenon, the simplest is the best” - Occam’s Razor

Variable Selection

- Backgrounds
 - ✓ Theoretically, model performance improves when the number of variables increases
(Under variable independence condition)
 - ✓ In reality, model performance degenerates due to variable dependence, existence of noise, etc.
- Purpose
 - ✓ Identify a subset of variables that best fit the model
- Effect
 - ✓ Remove correlations between variables
 - ✓ Simplified post-processing
 - ✓ Remove redundant or unnecessary variables while keeping relevant information
 - ✓ Visualization can be possible

Variable Selection

- Supervised vs. Unsupervised Dimensionality Reduction

- ✓ Supervised dimensionality reduction

- Use data mining models to verify the reduced dimensions
 - Dimensionality reduction results can be different according to the data mining algorithms employed

- ✓ Unsupervised dimensionality reduction

- Do not use data mining models during the process
 - Dimensionality reduction results are identical if the data and method is same

Variable Selection


- Dimensionality reduction techniques

- ✓ Variable/feature selection

- Select a subset of variables from the original variable set
 - Filter – Variable selection and model training are independent
 - Wrapper – Variable selection is done to optimize the result of the considered data mining model

- ✓ Variable/feature extraction

- Extract a new smaller set of variables that preserve the characteristics of the original data
 - Performance metric that is independent from data mining models is used



Variable Selection

- Selection vs. Extraction

✓ Conceptual difference between variable selection and variable extraction

X_1	X_2	X_3	...	X_n
...
...
...
...
...

Variable selection

X_1	X_5	X_8
...
...
...
...
...

Variable extraction

Z_1	Z_2	Z_3
...
...
...
...
...

$$Z_1 = X_1 + 0.2 * X_2$$

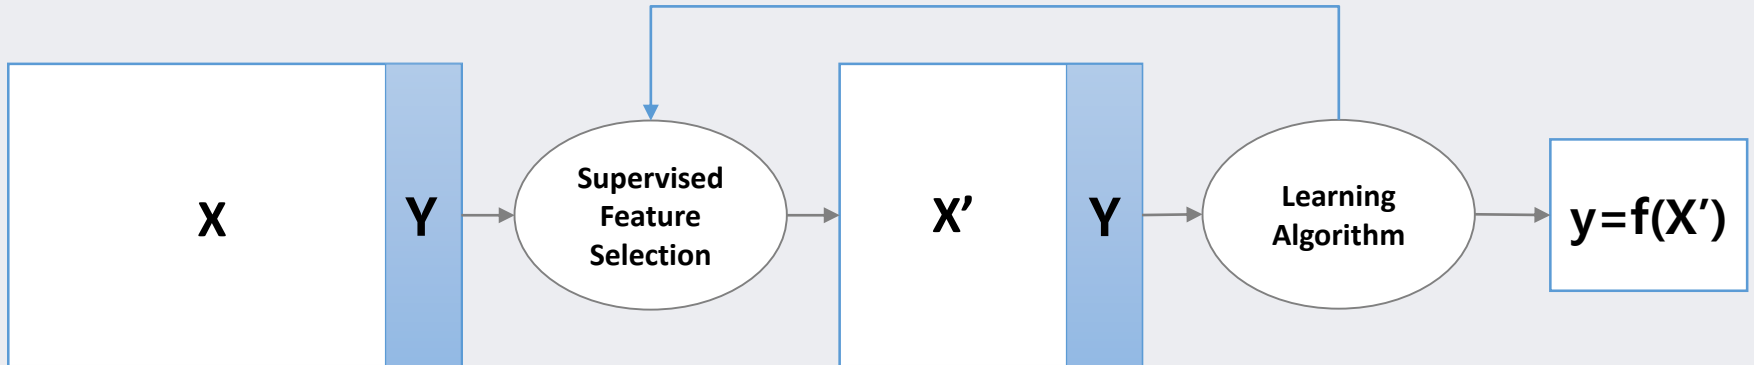
$$Z_2 = X_3 - 2 * X_5$$

$$Z_3 = X_4 + X_6 - X_9$$

Variable Selection

- Supervised variable selection

- ✓ Select d' variables from d variables ($d' \ll d$) in order to optimize the performance of the considered learning algorithm




- ✓ Select the learning algorithm before variable selection
- ✓ Different variable selection results are possible due to the variety of learning algorithms

Variable Selection


- Exhaustive search

- ✓ Search all possible combinations

- Ex) 3 variables



- A total of 6 possible subsets are tested




- ✓ Performance criteria for variable selection

- Akaike Information Criteria (AIC), Bayesian Information Criteria (BIC), Adjusted R^2 , Mallows's C_p , etc.

Variable Selection

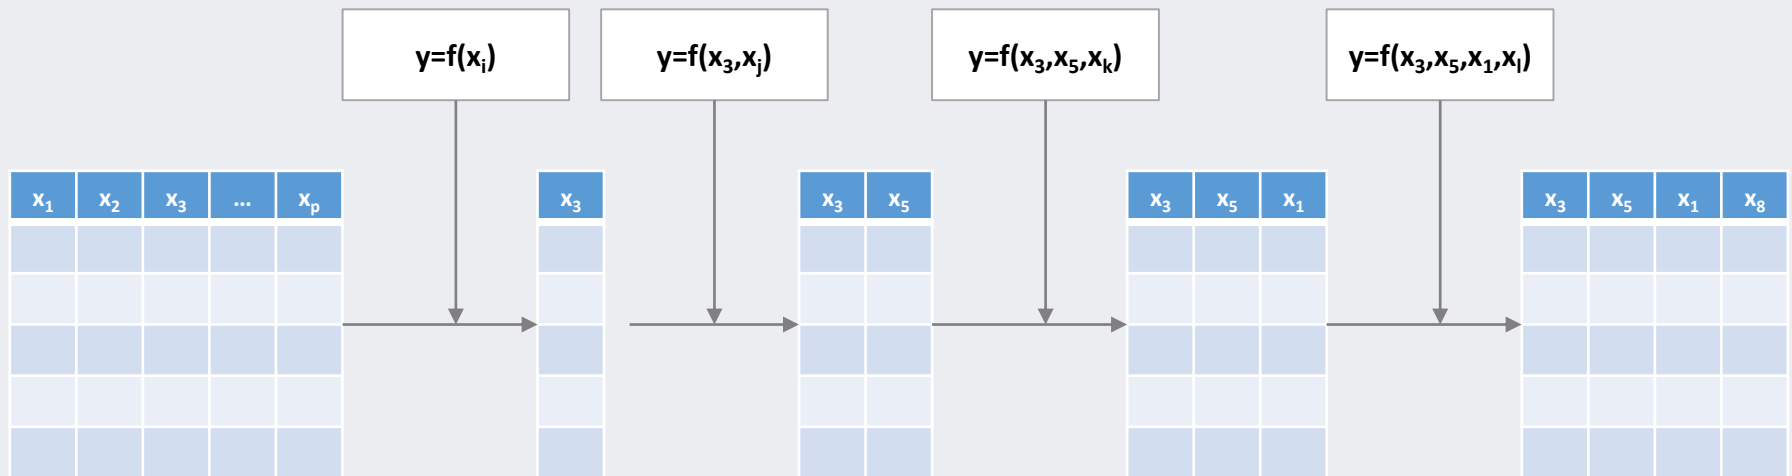
- Exhaustive search
 - ✓ Assume that we have a computer that can evaluate 10,000 models/second



Variable Selection

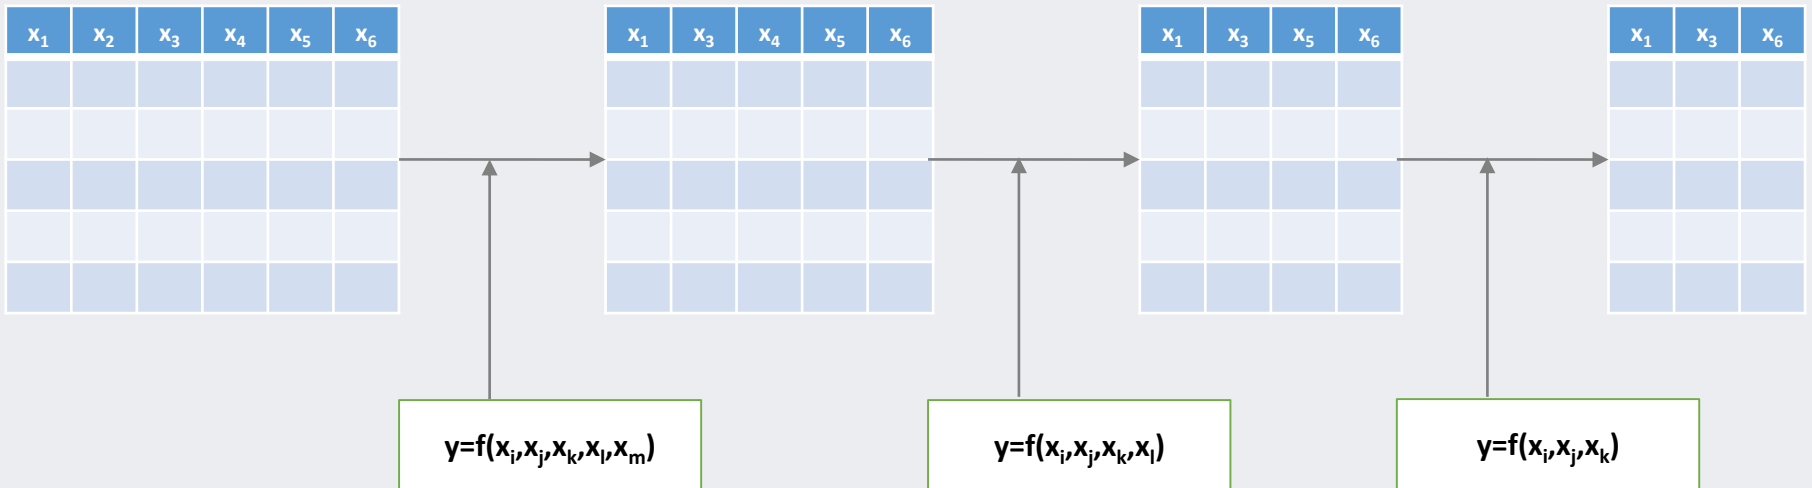
- Forward selection

- ✓ From the model with no variable, significant variables are sequentially added
- ✓ Once the variable is selected, it will never be removed (The number of variables gradually increases)



Variable Selection

- Backward Elimination
 - ✓ From the model with all variables, irrelevant variables are sequentially removed
 - ✓ Once a variable is removed, it will never be selected (The number of variables gradually decreases)



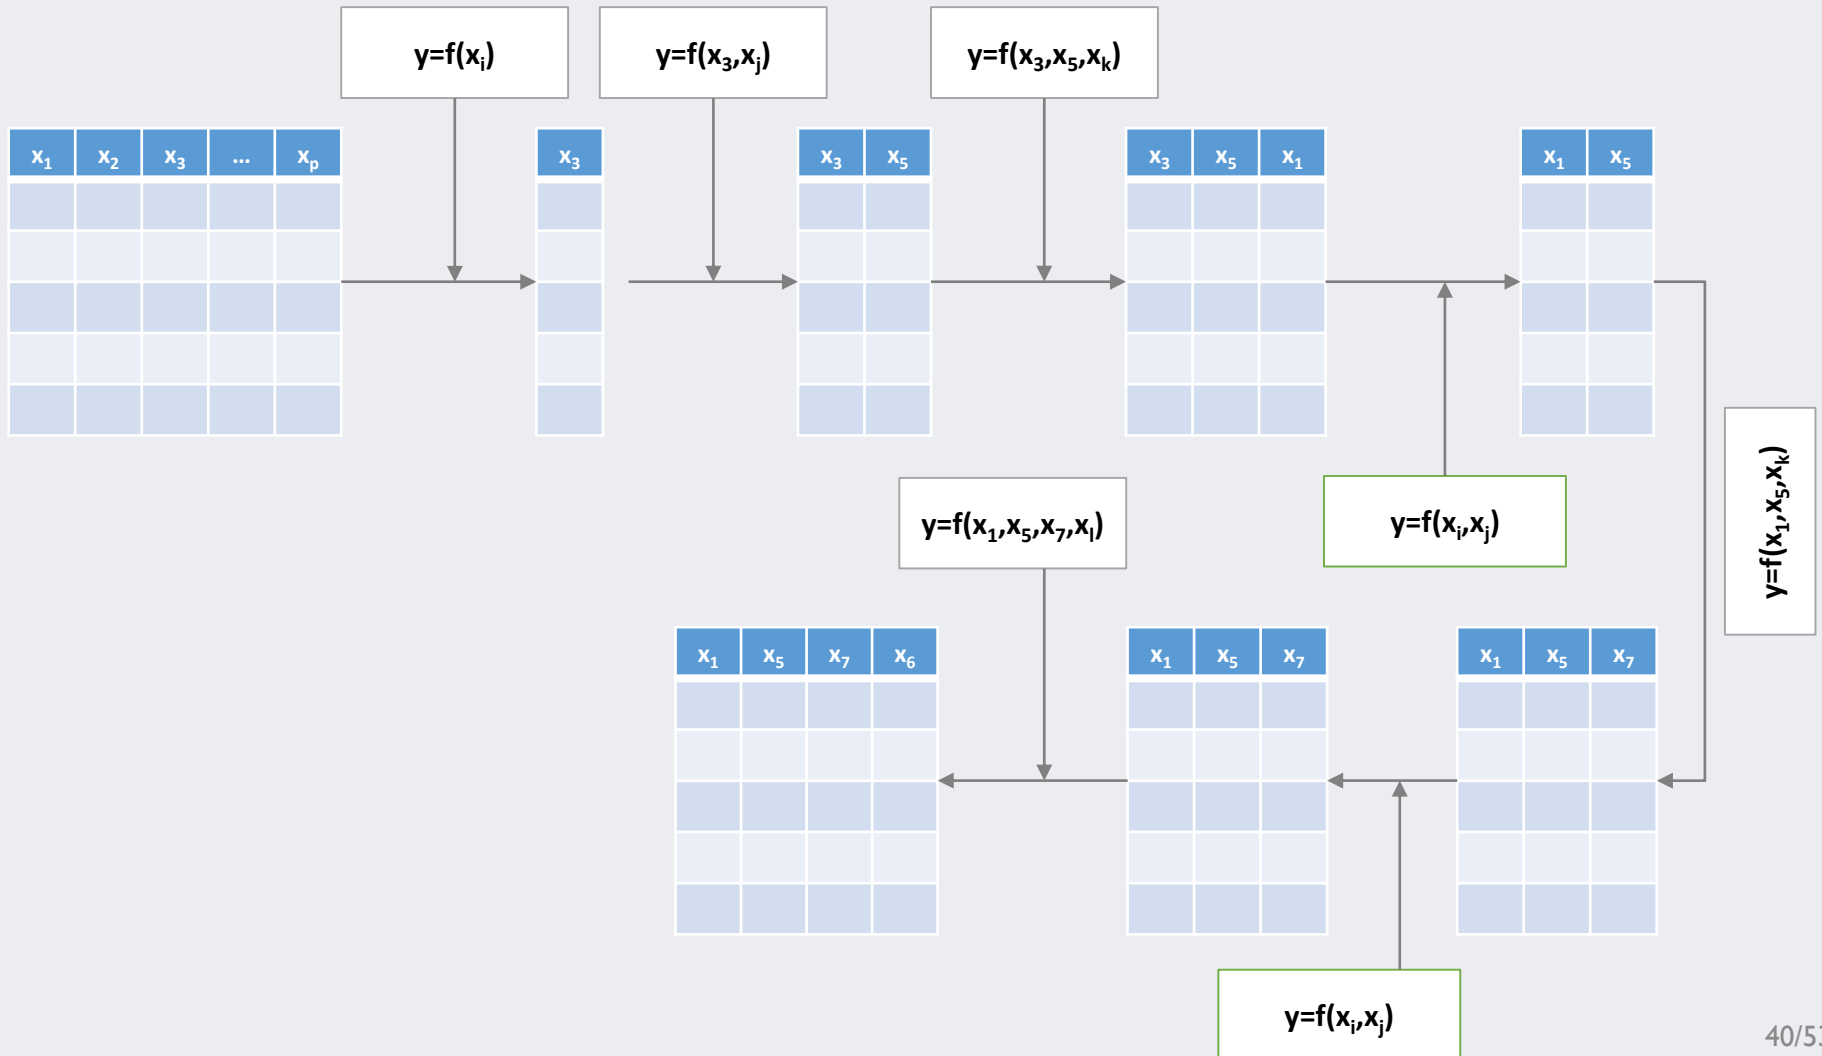
Variable Selection

- Stepwise Selection

- ✓ From the model with no variable, conduct the forward selection and backward elimination alternately
- ✓ Takes longer time than forward selection/backward elimination, but has more chances to find the optimal set of variables
- ✓ Variables that is either selected/removed can be reconsidered for selection/removal
- ✓ The number of variables increases in the early period, but it can either increase or decrease

Variable Selection

- Stepwise selection example



AGENDA

- 01 Multiple Linear Regression
- 02 Evaluating Regression Models
- 03 Variable Selection
- 04 **R Exercise**

R Exercise

- Data Set: Toyota Corolla Selling Price



Variable	Description	Variable	Description
		Guarantee_Period	Guarantee period in months
		ABS	Anti-Lock Brake System (Yes=1, No=0)
Price	Offer Price in EUROS	Airbag_1	Driver Airbag (Yes=1, No=0)
Age_08_04	Age in months as in August 2004	Airbag_2	Passenger Airbag (Yes=1, No=0)
Mfg_Month	Manufacturing month (1-12)	Airco	Airconditioning (Yes=1, No=0)
Mfg_Year	Manufacturing Year	Automatic_airco	Automatic Airconditioning (Yes=1, No=0)
KM	Accumulated Kilometers on odometer	Boardcomputer	Boardcomputer (Yes=1, No=0)
Fuel_Type	Fuel Type (Petrol, Diesel, CNG)	CD_Player	CD Player (Yes=1, No=0)
HP	Horse Power	Central_Lock	Central Lock (Yes=1, No=0)
Met_Color	Metallic Color? (Yes=1, No=0)	Powered_Windows	Powered Windows (Yes=1, No=0)
Automatic	Automatic (Yes=1, No=0)	Power_Steering	Power Steering (Yes=1, No=0)
CC	Cylinder Volume in cubic centimeters	Radio	Radio (Yes=1, No=0)
Doors	Number of doors	Mistlamps	Mistlamps (Yes=1, No=0)
Cylinders	Number of cylinders	Sport_Model	Sport Model (Yes=1, No=0)
Gears	Number of gear positions	Backseat_Divider	Backseat Divider (Yes=1, No=0)
Quarterly_Tax	Quarterly road tax in EUROS	Metallic_Rim	Metallic Rim (Yes=1, No=0)
Weight	Weight in Kilograms	Radio_cassette	Radio Cassette (Yes=1, No=0)
Mfr_Guarantee	Within Manufacturer's Guarantee period (Yes=1, No=0)	Parking_Assistant	Parking assistance system (Yes=1, No=0)
BOVAG_Guarantee	BOVAG (Dutch dealer network) Guarantee (Yes=1, No=0)	Tow_Bar	Tow Bar (Yes=1, No=0)

R Exercise

- Import the dataset & preprocessing

- ✓ Convert categorical variable to binary variables (1-of-c coding)

```

1 # Working directory 지정
2 setwd("C:\\RStudy")
3
4 # 실습 1: 전진선택/후진 소거/단계적선택 -----
5 # 분석에 필요한 패키지 설치 및 불러오기
6 # Multivariate linear regression
7 corolla <- read.csv("ToyotaCorolla.csv")
8
9 # Indices for the activated input variables
10 nCar <- dim(corolla)[1]
11 nVar <- dim(corolla)[2]
12
13 id_idx <- c(1,2)
14 category_idx <- 8
15
16 # 범주형 변수를 이진형 변수로 변환
17 dummy_p <- rep(0,nCar)
18 dummy_d <- rep(0,nCar)
19 dummy_c <- rep(0,nCar)
20
21 p_idx <- which(corolla$Fuel_Type == "Petrol")
22 d_idx <- which(corolla$Fuel_Type == "Diesel")
23 c_idx <- which(corolla$Fuel_Type == "CNG")
24
25 dummy_p[p_idx] <- 1
26 dummy_d[d_idx] <- 1
27 dummy_c[c_idx] <- 1
28
29 Fuel <- data.frame(dummy_p, dummy_d, dummy_c)
30 names(Fuel) <- c("Petrol", "Diesel", "CNG")
31
32 # Prepare the data for MLR
33 mlr_data <- cbind(corolla[, -c(id_idx, category_idx)], Fuel)

```

Price	Age_08_04	Mfg_Month	Mfg_Year	KM	Fuel_Type	HP	Met_Color	Automatic	cc
13500	23	10	2002	46986	Diesel	90	1	0	2000
13750	23	10	2002	72937	Diesel	90	1	0	2000
13950	24	9	2002	41711	Diesel	90	1	0	2000
14950	26	7	2002	48000	Diesel	90	0	0	2000
13750	30	3	2002	38500	Diesel	90	0	0	2000
12950	32	1	2002	61000	Diesel	90	0	0	2000
16900	27	6	2002	94612	Diesel	90	1	0	2000
18600	30	3	2002	75889	Diesel	90	1	0	2000
21500	27	6	2002	19700	Petrol	192	0	0	1800
12950	23	10	2002	71138	Diesel	69	0	0	1900
20950	25	8	2002	31461	Petrol	192	0	0	1800
19950	22	11	2002	43610	Petrol	192	0	0	1800
19600	25	8	2002	32189	Petrol	192	0	0	1800
21500	31	2	2002	23000	Petrol	192	1	0	1800
22500	32	1	2002	34131	Petrol	192	1	0	1800

KM	HP	Met_Color
46986	90	1
72937	90	1
41711	90	1
48000	90	0
38500	90	0
61000	90	0
94612	90	1
75889	90	1
19700	192	0
71138	69	0
31461	192	0
43610	192	0
32189	192	0
23000	192	1

...

Petrol	Diesel	CNG
0	1	0
0	1	0
0	1	0
0	1	0
0	1	0
0	1	0
0	1	0
0	1	0
1	0	0
0	1	0
1	0	0
1	0	0
1	0	0
1	0	0

R Exercise

- Divide the dataset into training/validation

```
35 # Split the data into the training/validation sets
36 trn_idx <- sample(1:nCar, round(0.7*nCar))
37 trn_data <- mlr_data[trn_idx,]
38 val_data <- mlr_data[-trn_idx,]
```

▶ mlr_data	1436 obs. of 37 variables
▶ trn_data	1005 obs. of 37 variables
▶ val_data	431 obs. of 37 variables

- MLR with all variables

```
40 # Train the MLR
41 full_model <- lm(Price ~ ., data = trn_data)
42 full_model
43 summary(full_model)
44 plot(full_model)
45
46 # Plot the result
47 plot(trn_data$Price, fitted(full_model), xlim = c(4000,35000), ylim = c(4000,35000))
48 abline(0,1,lty=3)
49
50 anova(full_model)
51 plot(fitted(full_model), resid(full_model), xlab="Fitted values", ylab="Residuals")
```

R Exercise

- MLR with all variables


```
> summary(full_model)

Call:
lm(formula = Price ~ ., data = trn_data)

Residuals:
    Min       1Q   Median       3Q      Max
-6571.9  -640.9   -49.0   624.2  5972.3

Coefficients: (3 not defined because of singularities)
              Estimate Std. Error t value Pr(>|t|)
(Intercept)   3.540e+03  1.724e+03   2.054  0.040257 *
Age_08_04     -1.178e+02  3.914e+00 -30.104 < 2e-16 ***
Mfg_Month     -1.059e+02  1.034e+01 -10.244 < 2e-16 ***
Mfg_Year              NA              NA      NA      NA
KM            -1.710e-02  1.338e-03 -12.777 < 2e-16 ***
HP             1.911e+01  3.601e+00   5.305  1.39e-07 ***
Met_Color     -4.358e+01  7.632e+01  -0.571  0.568134
Automatic      3.746e+02  1.458e+02   2.568  0.010368 *
cc            -5.613e-02  7.515e-02  -0.747  0.455279
Doors          7.198e+01  4.111e+01   1.751  0.080257 .
cylinders              NA              NA      NA      NA
Gears          1.959e+02  2.142e+02   0.915  0.360617
Quarterly_Tax  1.159e+01  2.128e+00   5.446  6.52e-08 ***
weight        8.879e+00  1.227e+00   7.233  9.54e-13 ***
Mfr_Guarantee  2.360e+02  7.381e+01   3.198  0.001430 **
BOVAG_Guarantee 3.989e+02  1.316e+02   3.033  0.002490 **
Guarantee_Period 7.207e+01  1.459e+01   4.938  9.27e-07 ***
ABS           -4.715e+01  1.300e+02  -0.363  0.716844
Airbag_1       4.498e+02  2.570e+02   1.750  0.080375 .
Airbag_2      -2.007e+02  1.314e+02  -1.527  0.127121
Airco          2.245e+02  8.919e+01   2.517  0.012008 *
Automatic_airco 2.435e+03  1.889e+02  12.890 < 2e-16 ***
Boardcomputer  -2.099e+02  1.194e+02  -1.758  0.078992 .
CD_Player      8.442e+01  1.010e+02   0.836  0.403239
Central_Lock   -7.471e+01  1.419e+02  -0.526  0.598678
Powered_windows 5.112e+02  1.424e+02   3.589  0.000349 ***
Power_steering -5.689e+02  2.842e+02  -2.002  0.045581 *
radio          5.575e+02  6.295e+02   0.886  0.376037
Mistlamps      1.869e+01  1.102e+02   0.170  0.865286
Sport_Model    2.790e+02  8.906e+01   3.132  0.001787 **
Backseat_Divider -6.961e+01  1.327e+02  -0.525  0.599953
Metallic_rim    5.536e+01  9.675e+01   0.572  0.567342
Radio_cassette  -5.593e+02  6.299e+02  -0.888  0.374863
Tow_Bar        -1.990e+02  8.018e+01  -2.482  0.013216 *
Petrol         1.096e+03  4.339e+02   2.527  0.011663 *
Diesel         5.269e+02  4.128e+02   1.276  0.202180
CNG              NA              NA      NA      NA
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1060 on 971 degrees of freedom
Multiple R-squared:  0.9127, Adjusted R-squared:  0.9098
F-statistic: 307.7 on 33 and 971 DF,  p-value: < 2.2e-16
```



R Exercise

- Variable selection I: Forward selection

✓ Starts with zero variable and adds the most significant variable at once

```
--  
53 # 변수선택 1: 전진선택법  
54 # Upperbound formula 만들기  
55 tmp_x <- paste(colnames(trn_data)[-1], collapse=" + ")  
56 tmp_xy <- paste("Price ~ ", tmp_x, collapse = "")  
57 tmp_xy  
58 as.formula(tmp_xy)  
59  
60 forward_model <- step(lm(Price ~ 1, data = trn_data),  
61                       scope = list(upper = as.formula(tmp_xy), lower = Price ~ 1), direction="forward", trace=1)  
62 summary(forward_model)  
63 anova(forward_model)  
64  
65 # 각 단계에서 선택된 변수 표시  
66 forward_model$anova$Step  
67 forward_model$anova$AIC  
68  
69 # 선택된 변수에 따른 AIC 감소분 표시  
70 plot(forward_model$anova$AIC, pch = 17, cex=2, main = "AIC Decrease (Forward Selection)", xlab = "Number of Steps", ylab = "AIC")  
71 text(forward_model$anova$AIC, forward_model$anova$Step, cex=1, pos=3, col="blue")
```

R Exercise

- Variable selection I: Forward selection

✓ Variable selection results (36 variables → 20 variables)

```
> summary(forward_model)
```

Call:

```
lm(formula = Price ~ Mfg_Year + Automatic_airco + Weight + KM +  
    Powered_windows + HP + Quarterly_Tax + Guarantee_Period +  
    BOVAG_Guarantee + Petrol + Mfr_Guarantee + Sport_Model +  
    Airco + Tow_Bar + Airbag_2 + Automatic + Boardcomputer +  
    Power_Steering + Airbag_1 + Doors, data = trn_data)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-6747.2	-653.8	-53.8	640.8	5908.7

Coefficients:


	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-2.807e+06	8.542e+04	-32.857	< 2e-16	***
Mfg_Year	1.402e+03	4.286e+01	32.718	< 2e-16	***
Automatic_airco	2.451e+03	1.746e+02	14.037	< 2e-16	***
Weight	9.233e+00	1.166e+00	7.918	6.50e-15	***
KM	-1.734e-02	1.309e-03	-13.252	< 2e-16	***
Powered_windows	4.650e+02	8.300e+01	5.602	2.74e-08	***
HP	1.819e+01	3.235e+00	5.625	2.42e-08	***
Quarterly_Tax	1.146e+01	2.013e+00	5.694	1.63e-08	***
Guarantee_Period	7.624e+01	1.377e+01	5.535	3.98e-08	***
BOVAG_Guarantee	4.078e+02	1.268e+02	3.216	0.001342	**
Petrol	6.593e+02	2.956e+02	2.231	0.025933	*
Mfr_Guarantee	2.263e+02	7.214e+01	3.137	0.001757	**
Sport_Model	2.811e+02	8.226e+01	3.417	0.000659	***
Airco	2.430e+02	8.500e+01	2.859	0.004334	**
Tow_Bar	-2.203e+02	7.742e+01	-2.846	0.004523	**
Airbag_2	-2.167e+02	9.707e+01	-2.232	0.025847	*
Automatic	3.395e+02	1.428e+02	2.378	0.017586	*
Boardcomputer	-1.929e+02	1.126e+02	-1.713	0.087054	.
Power_Steering	-6.486e+02	2.715e+02	-2.389	0.017075	*
Airbag_1	4.773e+02	2.521e+02	1.893	0.058598	.
Doors	5.867e+01	3.958e+01	1.482	0.138578	

signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1056 on 984 degrees of freedom

Multiple R-squared: 0.9121, Adjusted R-squared: 0.9103

F-statistic: 510.6 on 20 and 984 DF, p-value: < 2.2e-16



R Exercise

- Variable selection 2: Backward elimination

✓ Starts with all variable and removes the most insignificant variable at once

```
73 # 변수선택 2: 후진소거법
74 backward_model <- step(full_model, scope = list(upper = as.formula(tmp_xy), lower = Price ~ 1), direction="backward", trace=1)
75 summary(backward_model)
76 anova(backward_model)
77
78 # 각 단계에서 제거된 변수 표시
79 backward_model$anova$Step
80
81 # 제거된 변수에 따른 AIC 감소분 표시
82 plot(backward_model$anova$AIC, pch = 15, cex=2, main = "AIC Decrease (Backward Selection)", xlab = "Number of Steps", ylab = "AIC")
83 text(backward_model$anova$AIC, backward_model$anova$Step, cex=1, pos=3, col="red")
```

```
> backward_model$anova$step
[1] ""                "- CNG"            "- cylinders"       "- Mfg_Year"        "- Mistlamps"
[6] "- ABS"           "- Backseat_Divider" "- Central_Lock"    "- Met_Color"       "- Metallic_Rim"
[11] "- cc"            "- CD_Player"       "- Radio"           "- Radio_cassette"  "- Gears"
[16] "- Diesel"
```


R Exercise

- Variable selection 2: Backward elimination

✓ Variable selection results (36 variables → 21 variables)

```
> summary(backward_model)
```


```
Call:
lm(formula = Price ~ Age_08_04 + Mfg_Month + KM + HP + Automatic +
  Doors + Quarterly_Tax + Weight + Mfr_Guarantee + BOVAG_Guarantee +
  Guarantee_Period + Airbag_1 + Airbag_2 + Airco + Automatic_airco +
  Boardcomputer + Powered_windows + Power_Steering + Sport_Model +
  Tow_Bar + Petrol, data = trn_data)
```

```
Residuals:
    Min       1Q   Median       3Q      Max
-6744.2  -643.7   -43.5    630.5   5924.2
```

```
Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  4.528e+03  1.309e+03   3.460 0.000563 ***
Age_08_04    -1.176e+02  3.644e+00  -32.281 < 2e-16 ***
Mfg_Month    -1.067e+02  1.021e+01  -10.456 < 2e-16 ***
KM           -1.711e-02  1.327e-03  -12.898 < 2e-16 ***
HP            1.809e+01  3.236e+00   5.590 2.95e-08 ***
Automatic     3.373e+02  1.428e+02   2.363 0.018341 *
Doors         5.612e+01  3.965e+01   1.415 0.157290
Quarterly_Tax 1.156e+01  2.015e+00   5.738 1.27e-08 ***
Weight       9.259e+00  1.166e+00   7.938 5.57e-15 ***
Mfr_Guarantee 2.249e+02  7.215e+01   3.117 0.001879 **
BOVAG_Guarantee 4.138e+02  1.269e+02   3.260 0.001150 **
Guarantee_Period 7.511e+01  1.381e+01   5.437 6.82e-08 ***
Airbag_1     4.597e+02  2.526e+02   1.820 0.069067 .
Airbag_2    -2.272e+02  9.758e+01  -2.329 0.020071 *
Airco        2.377e+02  8.514e+01   2.791 0.005351 **
Automatic_airco 2.455e+03  1.746e+02  14.060 < 2e-16 ***
Boardcomputer -2.056e+02  1.133e+02  -1.816 0.069700 .
Powered_windows 4.620e+02  8.304e+01   5.563 3.41e-08 ***
Power_Steering -6.192e+02  2.729e+02  -2.269 0.023479 *
Sport_Model   2.717e+02  8.273e+01   3.284 0.001059 **
Tow_Bar      -2.156e+02  7.754e+01  -2.780 0.005531 **
Petrol        7.055e+02  2.988e+02   2.361 0.018404 *
```

```
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 1056 on 983 degrees of freedom
Multiple R-squared:  0.9122, Adjusted R-squared:  0.9103
F-statistic: 486.4 on 21 and 983 DF,  p-value: < 2.2e-16
```



R Exercise

- Variable selection 3: Stepwise selection

- ✓ Starts with zero variable and alternately adds the most significant variable and removes the most insignificant variable

```
85 # 변수선택 3: 단계적 선택법
86 stepwise_model <- step(lm(Price ~ 1, data = trn_data),
87   scope = list(upper = as.formula(tmp_xy), lower = Price ~ 1), direction="both", trace=1)
88 summary(stepwise_model)
89 anova(stepwise_model)
90
91 # 각 단계에서 선택/제거된 변수 표시
92 stepwise_model$anova$Step
93 stepwise_model$anova$AIC
94
95 # 제거/선택된 변수에 따른 AIC 감소분 표시
96 plot(stepwise_model$anova$AIC, pch = 19, cex=2, main = "AIC Decrease (Stepwise Selection)", xlab = "Number of Steps", ylab = "AIC")
97 text(stepwise_model$anova$AIC, stepwise_model$anova$Step, cex=1, pos=3, col="black")
```



```
> stepwise_model$anova$Step
[1] ""           "+ Mfg_Year"      "+ Automatic_airco" "+ Weight"         "+ KM"             "+ Powered_windows"
[7] "+ HP"        "+ Quarterly_Tax" "+ Guarantee_Period" "+ BOVAG_Guarantee" "+ Petrol"          "+ Mfr_Guarantee"
[13] "+ Sport_Model" "+ Airco"         "+ Tow_Bar"        "+ Airbag_2"       "+ Automatic"       "+ Boardcomputer"
[19] "+ Power_Steering" "+ Airbag_1"      "+ Doors"
```

R Exercise

- Variable selection 3: Stepwise selection

- ✓ Variable selection result

```
> summary(stepwise_model)
```

Call:

```
lm(formula = Price ~ Mfg_Year + Automatic_airco + Weight + KM +  
  Powered_windows + HP + Quarterly_Tax + Guarantee_Period +  
  BOVAG_Guarantee + Petrol + Mfr_Guarantee + Sport_Model +  
  Airco + Tow_Bar + Airbag_2 + Automatic + Boardcomputer +  
  Power_Steering + Airbag_1 + Doors, data = trn_data)
```

Residuals:


	Min	1Q	Median	3Q	Max
	-6747.2	-653.8	-53.8	640.8	5908.7

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-2.807e+06	8.542e+04	-32.857	< 2e-16	***
Mfg_Year	1.402e+03	4.286e+01	32.718	< 2e-16	***
Automatic_airco	2.451e+03	1.746e+02	14.037	< 2e-16	***
Weight	9.233e+00	1.166e+00	7.918	6.50e-15	***
KM	-1.734e-02	1.309e-03	-13.252	< 2e-16	***
Powered_windows	4.650e+02	8.300e+01	5.602	2.74e-08	***
HP	1.819e+01	3.235e+00	5.625	2.42e-08	***
Quarterly_Tax	1.146e+01	2.013e+00	5.694	1.63e-08	***
Guarantee_Period	7.624e+01	1.377e+01	5.535	3.98e-08	***
BOVAG_Guarantee	4.078e+02	1.268e+02	3.216	0.001342	**
Petrol	6.593e+02	2.956e+02	2.231	0.025933	*
Mfr_Guarantee	2.263e+02	7.214e+01	3.137	0.001757	**
Sport_Model	2.811e+02	8.226e+01	3.417	0.000659	***
Airco	2.430e+02	8.500e+01	2.859	0.004334	**
Tow_Bar	-2.203e+02	7.742e+01	-2.846	0.004523	**
Airbag_2	-2.167e+02	9.707e+01	-2.232	0.025847	*
Automatic	3.395e+02	1.428e+02	2.378	0.017586	*
Boardcomputer	-1.929e+02	1.126e+02	-1.713	0.087054	.
Power_Steering	-6.486e+02	2.715e+02	-2.389	0.017075	*
Airbag_1	4.773e+02	2.521e+02	1.893	0.058598	.
Doors	5.867e+01	3.958e+01	1.482	0.138578	

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1056 on 984 degrees of freedom
Multiple R-squared: 0.9121, Adjusted R-squared: 0.9103
F-statistic: 510.6 on 20 and 984 DF, p-value: < 2.2e-16



R Exercise

- Performance comparison among the variable selection techniques

```
99 # 검증 데이터에 대한 각 변수선택 결과의 예측 정확도 비교
100 full_haty <- predict(full_model, newdata = val_data)
101 forward_haty <- predict(forward_model, newdata = val_data)
102 backward_haty <- predict(backward_model, newdata = val_data)
103 stepwise_haty <- predict(stepwise_model, newdata = val_data)
104
105 # 회귀분석 예측성능 평가지표
106 # 1: Mean squared error (MSE)
107 perf_mat <- matrix(0,4,6)
108 perf_mat[1,1] <- mean((val_data$Price-full_haty)^2)
109 perf_mat[1,2] <- mean((val_data$Price-forward_haty)^2)
110 perf_mat[1,3] <- mean((val_data$Price-backward_haty)^2)
111 perf_mat[1,4] <- mean((val_data$Price-stepwise_haty)^2)
112
113 # 2: Root mean squared error (RMSE)
114 perf_mat[2,1] <- sqrt(mean((val_data$Price-full_haty)^2))
115 perf_mat[2,2] <- sqrt(mean((val_data$Price-forward_haty)^2))
116 perf_mat[2,3] <- sqrt(mean((val_data$Price-backward_haty)^2))
117 perf_mat[2,4] <- sqrt(mean((val_data$Price-stepwise_haty)^2))
118
119 # 3: Mean absolute error (MAE)
120 perf_mat[3,1] <- mean(abs(val_data$Price-full_haty))
121 perf_mat[3,2] <- mean(abs(val_data$Price-forward_haty))
122 perf_mat[3,3] <- mean(abs(val_data$Price-backward_haty))
123 perf_mat[3,4] <- mean(abs(val_data$Price-stepwise_haty))
124
125 # 4: Mean absolute percentage error (MAPE)
126 perf_mat[4,1] <- mean(abs((val_data$Price-full_haty)/val_data$Price))*100
127 perf_mat[4,2] <- mean(abs((val_data$Price-forward_haty)/val_data$Price))*100
128 perf_mat[4,3] <- mean(abs((val_data$Price-backward_haty)/val_data$Price))*100
129 perf_mat[4,4] <- mean(abs((val_data$Price-stepwise_haty)/val_data$Price))*100
130
131 # 변수선택 기법 결과 비교
132 rownames(perf_mat) <- c("MSE", "RMSE", "MAE", "MAPE")
133 colnames(perf_mat) <- c("All", "Forward", "Backward", "Stepwise", "GA_default", "GA_yourOwn")
134 perf_mat
```

```
> perf_mat
```

	All	Forward	Backward	Stepwise
MSE	1.577365e+06	1.634343e+06	1.623485e+06	1.634343e+06
RMSE	1.255932e+03	1.278414e+03	1.274160e+03	1.278414e+03
MAE	9.121387e+02	9.242534e+02	9.211011e+02	9.242534e+02
MAPE	9.428209e+00	9.538071e+00	9.497384e+00	9.538071e+00

