



A man is known by the  
company he keeps.

Aesop

“quote fancy”

# Lecture 7: k-Nearest Neighbor

Pilsung Kang

School of Industrial Management Engineering

Korea University

# AGENDA

**01** k-Nearest Neighbor Classification

---

**02** k-Nearest Neighbor Regression

---

**03** R Exercise

---

# k-Nearest Neighbor Classification

- Gender classification



Men


Vs.

Women




# k-Nearest Neighbor Classification

- Which class does the question mark belong to?




# k-Nearest Neighbor Classification

- Which class does the question mark belong to?



# k-Nearest Neighbor Classification

- Which class does the question mark belong to?



# k-Nearest Neighbor Classification

- Motivation

類類相從 近墨者黑

“Birds of a feather flock together”

“A Man is known by the company he keeps”

# k-Nearest Neighbor Classification

## k-NN Classification Process

- Step I: Prepare the reference data
  - ✓ Define attributes
    - Height, Weight, Body Fat Statistics (BFS)
  - ✓ Collect sufficient number of records from each class

No.	Height	Weight	BFS	Class (Gender)
1	187	93	15	M
2	165	51	25	F
3	174	68	14	M
4	156	48	29	F
...	...	...	...	...
N	168	59	12	M




# k-Nearest Neighbor Classification

## k-NN Classification Process

- Step 2: Define the similarity measure
  - ✓ Similarity  $\propto 1/\text{distance}$
  - ✓ Minkovski distance with order p

$$\text{distance}(P = (x_1, x_2, \dots, x_n), Q = (y_1, y_2, \dots, y_n)) = \left( \sum_{i=1}^n |x_i - y_i|^p \right)^{\frac{1}{p}}$$



- ✓  $p=2$ : Euclidean distance
- ✓  $p=1$ : Manhattan distance

# k-Nearest Neighbor Classification

## k-NN Classification Process


- Step 3: Initialize the set of candidate values for  $k$ 
  - ✓ If  $k$  is too **small**, then the classification will be highly locally sensitive (**over-fitting**).
  - ✓ If  $k$  is too **large**, then it will lose the ability to capture the local structure (**under-fitting**).
  - ✓ A proper  $k$  should be chosen among a set of candidates.
  - ✓ Use the validation data.

# k-Nearest Neighbor Classification


## k-NN Classification Process

- Step 3: Initialize the set of candidate values for k


1-nearest neighbour




10-nearest neighbour



20-nearest neighbour



50-nearest neighbour



# k-Nearest Neighbor Classification

## k-NN Classification Process

- Step 4: Determine the combining rule

✓ Majority voting vs. Weighted voting

For a new data  <b>X</b>	Neighbor	Class	Distance	1/distance	Weight
	N1	M	1	1.00	0.44
	N2	F	2	0.50	0.22
	N3	M	3	0.33	0.15
	N4	F	4	0.25	0.11
	N5	F	5	0.20	0.08

✓ Majority voting:  $P(X=M) = 2/5 = 0.4$

✓ Weighted voting:  $P(X=M) = 0.59$

✓ If the cut-off is set to 0.5 X is classified as **F** by the majority voting while classified as **M** by the weighted voting

# k-Nearest Neighbor Classification

## k-NN Classification Process

- Step 5: Find the best k using the validation dataset

Value of k	% Error Training	% Error Validation
1	0.00	33.33
2	16.67	33.33
3	11.11	33.33
4	22.22	33.33
5	11.11	33.33
6	27.78	33.33
7	22.22	33.33
8	22.22	16.67
9	22.22	16.67
10	22.22	16.67
11	16.67	33.33
12	16.67	16.67
13	11.11	33.33
14	11.11	16.67
15	5.56	33.33
16	16.67	33.33
17	11.11	33.33
18	50.00	50.00

<--- Best k

# k-Nearest Neighbor Classification

- k-NN Issue 1: Normalization

- ✓ Normalization or scaling must be done before finding k-nearest neighbors
- ✓ If not, variables with large measuring units are over-emphasized while variables with small measuring units are under-evaluated

[Before Normalization]

No.	Height	Weight	BFS	Gender
1	187	93	15	M
2	165	51	25	F
3	174	68	14	M
4	156	48	29	F
...	...	...	...	...
N	168	59	12	M
Avg.	165	65	20	-
Stdev.	15	10	5	-

[After Normalization]

No.	Height	Weight	BFS	Gender
1	1.47	2.80	-1.00	M
2	0.00	-1.40	1.00	F
3	0.60	0.30	-1.20	M
4	-0.60	-1.70	1.80	F
...	...	...	...	...
N	0.20	-0.60	-1.60	M

# k-Nearest Neighbor Classification

- k-NN Issue 2: Cut-off
  - ✓ Consider the prior probability of each class
  - ✓ Assume that  $N(CM) = 100$ ,  $N(CF) = 400$

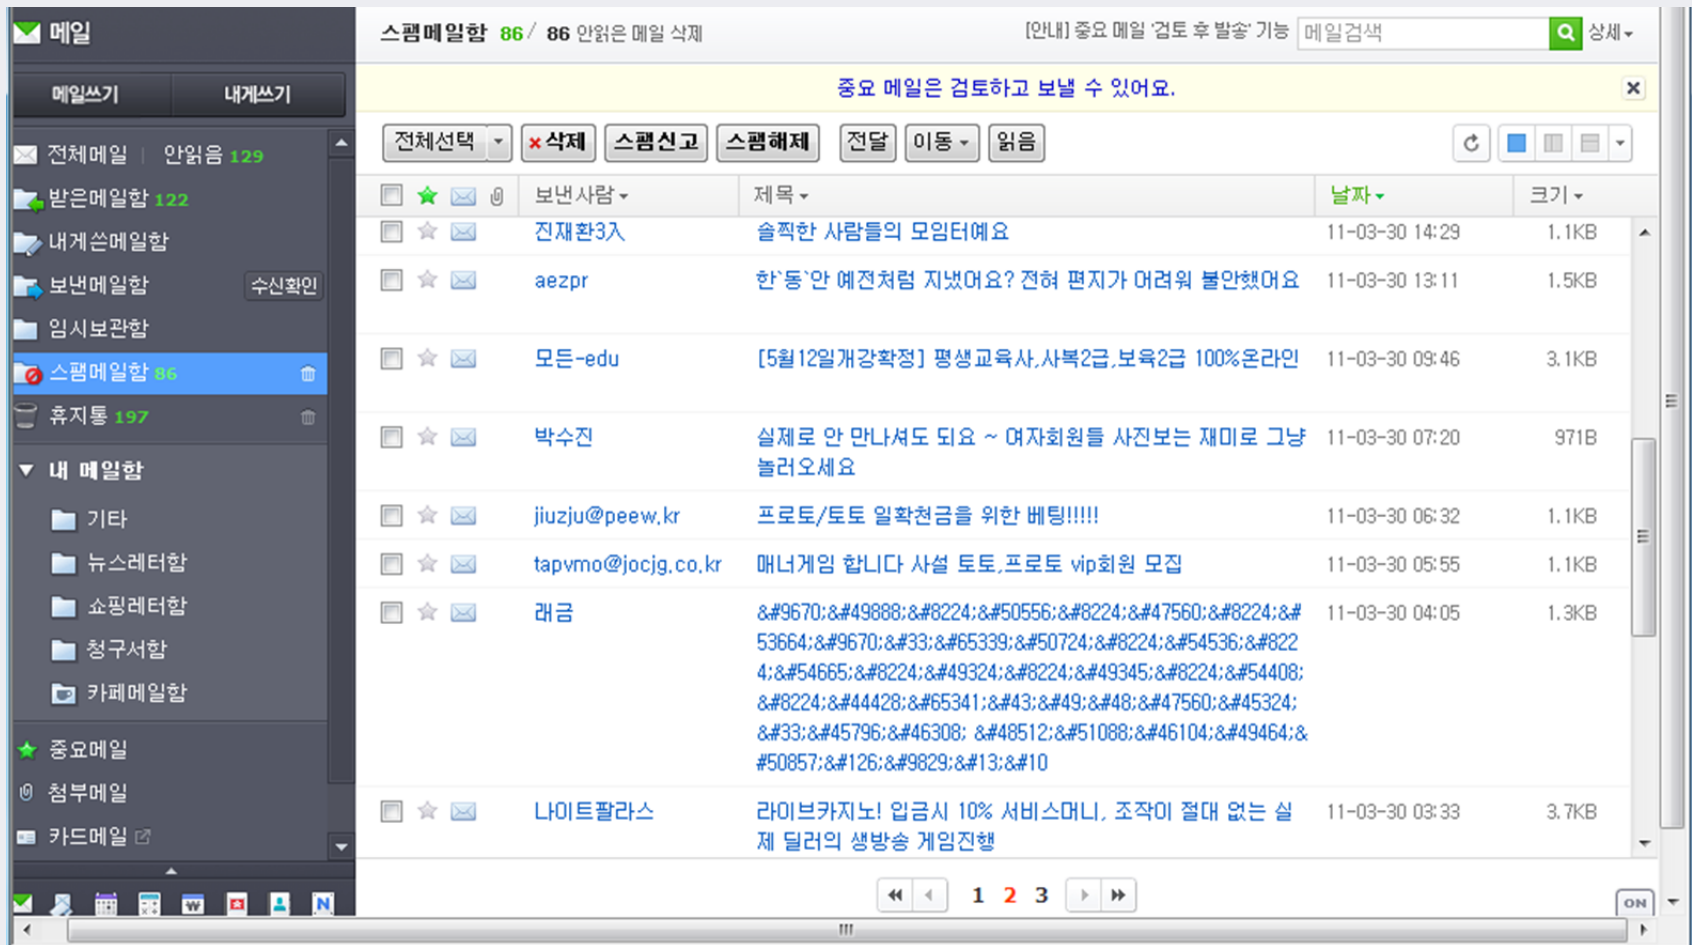
For a new data  <b>X</b>	Neighbor	Class
	N1	M
	N2	F
	N3	M
	N4	F
	N5	F

Majority voting  
 $P(X=M)=0.4$

- ✓ If the cut-off is set to 0.5 (assuming equal class distribution), then X is classified as **F**.
- ✓ If the cut-off is set to 0.2 (proportion of M among the people), then X is classified as **M**.

# k-Nearest Neighbor Classification

- k-NN Classification Application: Spam Filtering





# k-Nearest Neighbor Classification

- k-NN Classification Application: Spam Filtering


✓ Attributes: frequency of a set of keywords

Mail	회의	수정	기안	보고	대박	머니	외로워	미팅	...	스팸?
1	2	3	1	0	0	0	0	0	...	N
2	1	0	2	3	0	0	1	0	...	N
3	2	2	3	1	0	0	0	1	...	N
4	0	0	0	0	3	2	0	0	...	Y
5	0	0	0	1	0	0	2	3	...	Y
...	...	...	...	...	...	...	...	...	...	...

# k-Nearest Neighbor Classification

- k-NN Classification Application: Spam Filtering

- ✓ For a given new mail, find 5 similar existing mails



If we use the majority voting, then classify the mail as “spam”

If we use the weighted voting, then classify the mail as “no spam”

# AGENDA

**01** k-Nearest Neighbor Classification

---

**02** k-Nearest Neighbor Regression

---

**03** R Exercise

---

# k-Nearest Neighbor Regression

- Regression problem revisited



Predict  
One's  
BFS



10.0

21.7

8.9

19.9

23.4



28.9

15.7

21.6

21.5

23.2

# k-Nearest Neighbor Regression

## k-NN Regression Process

- Step 1: Prepare the reference data
  - ✓ Define attributes
    - Height, Weight, Gender
  - ✓ Collect sufficient number of records from each class


No.	Height	Weight	Gender (M=1)	BFS
1	187	93	1	15
2	165	51	0	25
3	174	68	1	14
4	156	48	0	29
...	...	...	...	...
N	168	59	1	12

# k-Nearest Neighbor Regression

## k-NN Regression Process

- Step 2: Define the similarity measure
  - ✓ Similarity  $\propto 1/\text{distance}$
  - ✓ Minkovski distance with order p

$$\text{distance}(P = (x_1, x_2, \dots, x_n), Q = (y_1, y_2, \dots, y_n)) = \left( \sum_{i=1}^n |x_i - y_i|^p \right)^{\frac{1}{p}}$$



- ✓ p=2: Euclidean distance
- ✓ p=1: Mahattan distance

# k-Nearest Neighbor Regression

## k-NN Regression Process

- Step 3: Initialize the set of candidate values for  $k$ 
  - ✓ If  $k$  is too **small**, then the classification will be highly locally sensitive (**over-fitting**).
  - ✓ If  $k$  is too **large**, then it will lose the ability to capture the local structure (**under-fitting**).
  - ✓ A proper  $k$  should be chosen among a set of candidates.
  - ✓ Use the validation data.

# k-Nearest Neighbor Regression

## k-NN Regression Process

- Step 4: Determine the combining rule

✓ Simple average vs. Weighted average

For a  
new  
data

**X**

Neighbor	BFS	Distance	1/distance	Weight
N1	15.4	1	1.00	0.44
N2	17.2	2	0.50	0.22
N3	12.3	3	0.33	0.15
N4	11.5	4	0.25	0.11
N5	10.9	5	0.20	0.08

✓ Simple average

- $\text{BFS of } X = (15.4 + 17.2 + 12.3 + 11.5 + 10.9) / 5 = 13.46$

✓ Weighted average

- $\text{BFS of } X = 0.44 * 15.4 + 0.22 * 17.2 + 0.15 * 12.3 + 0.11 * 11.5 + 0.08 * 10.9 = 14.54$



# k-Nearest Neighbor Regression

## k-NN Regression Process

- Step 5: Find the best k using the validation dataset

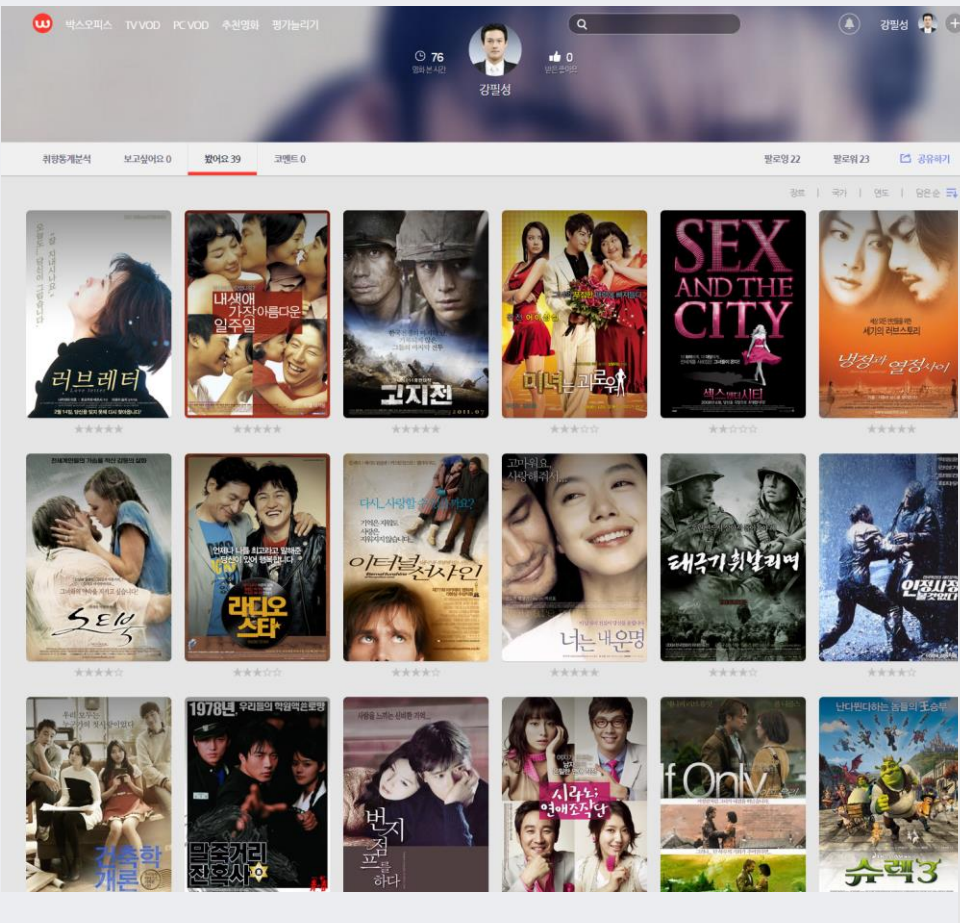
Value of k	% Error Training	% Error Validation
1	0.00	33.33
2	16.67	33.33
3	11.11	33.33
4	22.22	33.33
5	11.11	33.33
6	27.78	33.33
7	22.22	33.33
8	22.22	16.67
9	22.22	16.67
10	22.22	16.67
11	16.67	33.33
12	16.67	16.67
13	11.11	33.33
14	11.11	16.67
15	5.56	33.33
16	16.67	33.33
17	11.11	33.33
18	50.00	50.00

<--- Best k

# k-Nearest Neighbor Regression

- k-Nearest Neighbor Regression Application

- ✓ Collaborative Filtering-based Recommendation



# k-Nearest Neighbor Regression

- Recommendation List

예상 별점이 가장 높은 영화 ×

장르 | 국가 | 연도 | 추천이유

<p>4.3</p> <p>예상별점 4.3개 ★★★★★</p> <p>&lt;내 생애 가장 아름다운 일주일&gt;과 비슷해요.</p>	<p>4.2</p> <p>&lt;올의 크리스마스&gt;와 비슷해요.</p>	<p>4.0</p> <p>&lt;시리노 연애조각단&gt;과 비슷해요.</p>	<p>4.3</p> <p>좋아하는 배우 설경구</p>
<p>4.2</p> <p>&lt;친구&gt;와 비슷해요.</p>	<p>4.0</p> <p>&lt;너는 내 운명&gt;과 비슷해요.</p>	<p>4.0</p> <p>예상별점 4.0개 ★★★★★</p>	<p>4.3</p> <p>예상별점 4.3개 ★★★★★</p>
<p>4.6</p> <p>&lt;이탈리아 선사건&gt;의 2편과 비슷해요.</p>	<p>4.0</p> <p>&lt;우리들의 행복한 시간&gt;과 비슷해요.</p>	<p>4.0</p> <p>&lt;노트북&gt;의 1편과 비슷해요.</p>	<p>4.0</p> <p>예상별점 4.0개 ★★★★★</p>

# k-Nearest Neighbor Regression

- Collaborative Filtering

	Movie 1	Movie 2	Movie 3	Movie 4	Movie 5	...	Movie N
PSKang	10	9	5	6	9	...	? <b>9</b>

Cust.	Movie 1	Movie 2	Movie 3	Movie 4	Movie 5	...	Movie N
<b>1</b>	10	8	4	7	10	...	10
2	8	5	7	9	4	...	5
<b>3</b>	10	9	6	5	8	...	9
4	4	2	10	10	5	...	3
5	7	4	6	8	5	...	3
6	5	2	10	10	10	...	6
<b>7</b>	10	8	6	6	8	...	8
...	...	...	...	...	...	...	...
N	5	7	1	5	4	...	7

# AGENDA

**01** k-Nearest Neighbor Classification

---

**02** k-Nearest Neighbor Regression

---


**03** R Exercise

---



# R Exercise

- k-NN Illustration



# R Exercise

- Illustration with 2-D synthetic data


```
# k-Nearest Neighbor Illustration -----
install.packages("ElemStatLearn", dependencies = TRUE)
install.packages("class", dependencies = TRUE)
library(ElemStatLearn)
library(class)

# 2-D artificial data example with k=1
x <- mixture.example$x
g <- mixture.example$y
xnew <- mixture.example$xnew
mod1 <- knn(x, xnew, g, k=50, prob=TRUE)
prob1 <- attr(mod1, "prob")
prob1 <- ifelse(mod1=="1", prob1, 1-prob1)
px1 <- mixture.example$px1
px2 <- mixture.example$px2
prob1 <- matrix(prob1, length(px1), length(px2))
par(mar=rep(2,4))
contour(px1, px2, prob1, levels=0.5, labels="", xlab="", ylab="", main= "50-nearest neighbour", axes=FALSE)
points(x, col=ifelse(g==1, "coral", "cornflowerblue"))
gd <- expand.grid(x=px1, y=px2)
points(gd, pch=".", cex=1.2, col=ifelse(prob1>0.5, "coral", "cornflowerblue"))
box()
```

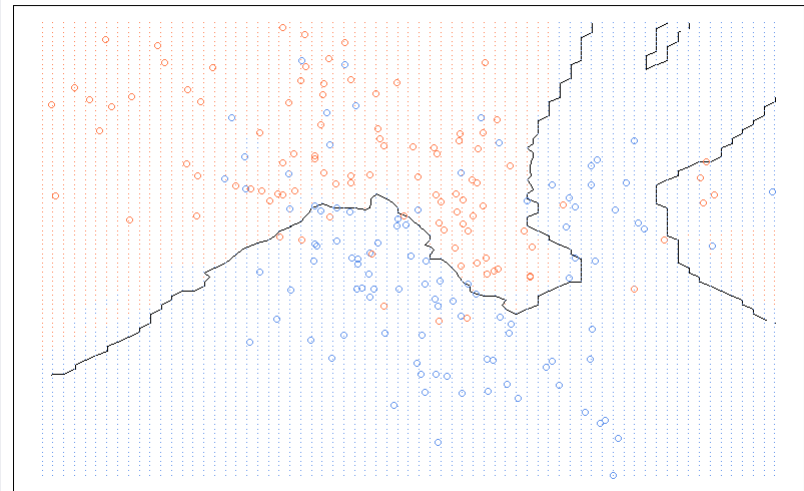
# R Exercise

- Decision boundaries with regard to different k values


1-nearest neighbour




10-nearest neighbour



20-nearest neighbour



50-nearest neighbour





# R Exercise

- k-NN Classification & Regression
  - ✓ Classification: package “kkn”
  - ✓ Regression: package “FNN”
- Classification dataset: Wisconsin breast cancer data
  - ✓ 569 patients (malignant: 357, benign: 212)
  - ✓ To determine whether the current patient is malignant or benign based on the 30 variables extracted from the 3D image of cell nucleus

Ten real-valued features are computed for each cell nucleus:

- a) radius (mean of distances from center to points on the perimeter)
- b) texture (standard deviation of gray-scale values)
- c) perimeter
- d) area
- e) smoothness (local variation in radius lengths)
- f) compactness ( $\text{perimeter}^2 / \text{area} - 1.0$ )
- g) concavity (severity of concave portions of the contour)
- h) concave points (number of concave portions of the contour)
- i) symmetry
- j) fractal dimension (“coastline approximation” - 1)

# R Exercise

- k-NN Classification

- ✓ Divide the entire dataset into training (70%) and validation (30%) and evaluate the performance of the classifier with  $k=1$

- Install the package and load the dataset

```
25 # k-Nearest Neighbor Learning (Classification) -----
26 # kknn package install & call
27 install.packages("kknn", dependencies = TRUE)
28 library(kknn)
29
30 # Load the wdbc data
31 RawData <- read.csv("wdbc.csv", header = FALSE)
32 head(RawData)
```

```
> head(RawData)
```

	V1	V2	V3	V4	V5	V6	V7	V8	V9	V10	V11	V12	V13	V14	V15	V16
1	13.540	14.36	87.46	566.3	0.09779	0.08129	0.06664	0.047810	0.1885	0.05766	0.2699	0.7886	2.058	23.560	0.008462	0.014600
2	13.080	15.71	85.63	520.0	0.10750	0.12700	0.04568	0.031100	0.1967	0.06811	0.1852	0.7477	1.383	14.670	0.004097	0.018980
3	9.504	12.44	60.34	273.9	0.10240	0.06492	0.02956	0.020760	0.1815	0.06905	0.2773	0.9768	1.909	15.700	0.009606	0.014320
4	13.030	18.42	82.61	523.8	0.08983	0.03766	0.02562	0.029230	0.1467	0.05863	0.1839	2.3420	1.170	14.160	0.004352	0.004899
5	8.196	16.84	51.71	201.9	0.08600	0.05943	0.01588	0.005917	0.1769	0.06503	0.1563	0.9567	1.094	8.205	0.008968	0.016460
6	12.050	14.63	78.04	449.3	0.10310	0.09092	0.06592	0.027490	0.1675	0.06043	0.2636	0.7294	1.848	19.870	0.005488	0.014270

	V17	V18	V19	V20	V21	V22	V23	V24	V25	V26	V27	V28	V29	V30	V31
1	0.02387	0.013150	0.01980	0.002300	15.110	19.26	99.70	711.2	0.14400	0.17730	0.23900	0.12880	0.2977	0.07259	M
2	0.01698	0.006490	0.01678	0.002425	14.500	20.49	96.09	630.5	0.13120	0.27760	0.18900	0.07283	0.3184	0.08183	M
3	0.01985	0.014210	0.02027	0.002968	10.230	15.66	65.13	314.9	0.13240	0.11480	0.08867	0.06227	0.2450	0.07773	M
4	0.01343	0.011640	0.02671	0.001777	13.300	22.81	84.46	545.9	0.09701	0.04619	0.04833	0.05013	0.1987	0.06169	M
5	0.01588	0.005917	0.02574	0.002582	8.964	21.96	57.26	242.2	0.12970	0.13570	0.06880	0.02564	0.3105	0.07409	M
6	0.02322	0.005660	0.01428	0.002422	13.760	20.70	89.88	582.6	0.14940	0.21560	0.30500	0.06548	0.2747	0.08301	M

# R Exercise

- k-NN Classification

- ✓ Divide the entire dataset into training (70%) and validation (30%) and evaluate the performance of the classifier with  $k=1$ 
  - Normalize the variables and divide the dataset

```
# Divide the dataset into the training (70%) and Validation (30%) datasets
trn_idx <- sample(1:length(Class), round(0.7*length(Class)))
trnInputs <- ScaledInputData[trn_idx,]
trnTargets <- Class[trn_idx]
valInputs <- ScaledInputData[-trn_idx,]
valTargets <- Class[-trn_idx]

trnData <- data.frame(trnInputs, trnTargets)
colnames(trnData)[31] <- "Target"
valData <- data.frame(valInputs, valTargets)
colnames(valData)[31] <- "Target"
```

▶ trnData	398 obs. of 31 variables
trnInputs	num [1:398, 1:30] 0.577 0.233 -0.658 -0.865 -0.484 ...
▶ valData	171 obs. of 31 variables
valInputs	num [1:171, 1:30] -0.138 -1.123 -1.447 -0.169 -0.36 ...

# R Exercise

- k-NN Classification

- ✓ Divide the entire dataset into training (70%) and validation (30%) and evaluate the performance of the classifier with  $k=1$

- Perform k-NN

```
# Perform k-nn classification with k=1, Distance = Euclidean, and weighted scheme = majority voting
kknn <- kknn(Target ~ ., trnData, valData, k=1, distance=2, kernel = "rectangular")

# View the k-nn results
summary(kknn)
kknn$CL
kknn$W
kknn$D
```

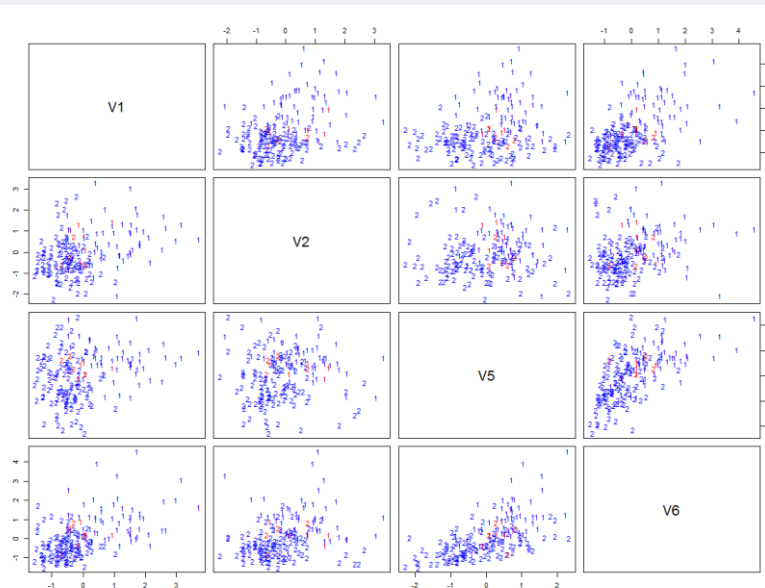
# R Exercise

- k-NN Classification

- ✓ Divide the entire dataset into training (70%) and validation (30%) and evaluate the performance of the classifier with  $k=1$

- Visualize the result

```
# Visualize the classification results
knnfit <- fitted(kknn)
table(valTargets, knnfit)
pcol <- as.character(as.numeric(valTargets))
pairs(valData[c(1,2,5,6)], pch = pcol, col = c("blue", "red")[valTargets != knnfit+1])
```



# R Exercise

- k-NN Classification

- ✓ Divide the entire dataset into training (70%) and validation (30%) and evaluate the performance of the classifier with  $k=1$

- Check the validation error

```
table(valTargets, kknn$fitted.values)
```

```
> table(valTargets, kknn$fitted.values)
```

```
valTargets    B    M
           B  57    5
           M   4 105
```

# R Exercise

- k-NN Classification

- ✓ Find the best k based on leave-one-out (LOO) validation

```
# Leave-one-out validation for finding the best k
knnttr <- train.kknn(Target ~ ., trnData, kmax=10, distance=2, kernel="rectangular")

knnttr$MISCLASS
knnttr$best.parameters
```

```
> knnttr$MISCLASS
rectangular
1  0.05778894
2  0.06783920
3  0.03517588
4  0.03517588
5  0.04020101
6  0.03015075
7  0.03266332
8  0.03517588
9  0.03768844
10 0.04020101
```

```
> knnttr$best.parameters
$kernel
[1] "rectangular"

$k
[1] 6
```

# R Exercise

- k-NN Classification

✓ Train the model with the best k

```
# Perform k-nn classification with the best k, Distance = Euclidean, and weighted scheme = majority voting
kkn_opt <- kkn(Target ~ ., trnData, valData, k=knntr$best.parameters$k, distance=2, kernel = "rectangular")
fit_opt <- fitted(kkn_opt)
cfmatrix <- table(valTargets, fit_opt)
cfmatrix
```

```
> cfmatrix
      fit_opt
valTargets B  M
      B  59  3
      M   2 107
```



# R Exercise

- k-NN Classification

✓ Evaluate the classification performance

```
# Summarize the classification performances
Cperf <- matrix(0, nrow=1, ncol=3)
colnames(Cperf) <- c("Accuracy", "BCR", "F1")

# Simple Accuracy
Cperf[1,1] <- (cfmatrix[1,1]+cfmatrix[2,2])/sum(cfmatrix)

# Balanced correction rate (BCR)
Cperf[1,2] <- sqrt((cfmatrix[1,1]/(cfmatrix[1,1]+cfmatrix[1,2]))*(cfmatrix[2,2]/(cfmatrix[2,1]+cfmatrix[2,2])))



# F1-measure
Recall <- cfmatrix[2,2]/(cfmatrix[2,1]+cfmatrix[2,2])
Precision <- cfmatrix[1,1]/(cfmatrix[1,1]+cfmatrix[1,2])
Cperf[1,3] <- 2*Recall*Precision/(Recall+Precision)

Cperf
```

```
> Cperf
      Accuracy      BCR      F1
[1,] 0.9707602 0.9665155 0.9663988
```

# R Exercise

- k-NN Regression
  - ✓ Dataset: Concrete Strength

		<u>Concrete Compressive Strength</u>	Multivariate	Regression	Real	1030	9	2007
---	---	--------------------------------------	--------------	------------	------	------	---	------

## Name -- Data Type -- Measurement -- Description

Cement (component 1) -- quantitative -- kg in a m3 mixture -- Input Variable  
Blast Furnace Slag (component 2) -- quantitative -- kg in a m3 mixture -- Input Variable  
Fly Ash (component 3) -- quantitative -- kg in a m3 mixture -- Input Variable  
Water (component 4) -- quantitative -- kg in a m3 mixture -- Input Variable  
Superplasticizer (component 5) -- quantitative -- kg in a m3 mixture -- Input Variable  
Coarse Aggregate (component 6) -- quantitative -- kg in a m3 mixture -- Input Variable  
Fine Aggregate (component 7) -- quantitative -- kg in a m3 mixture -- Input Variable  
Age -- quantitative -- Day (1~365) -- Input Variable  
Concrete compressive strength -- quantitative -- MPa -- Output Variable

# R Exercise

- k-NN Regression

- ✓ Install the package and normalize the data

```
# k-Nearest Neighbor Learning (Regression) -----  
install.packages("FNN", dependencies = TRUE)  
library(FNN)  
# Concrete strength data  
concrete <- read.csv("concrete.csv")  
  
RegX <- concrete[,1:8]  
RegY <- concrete[,9]  
  
# Data Normalization  
RegX <- scale(RegX, center = TRUE, scale = TRUE)  
  
# Combine X and Y  
RegData <- as.data.frame(cbind(RegX, RegY))
```

# R Exercise

- k-NN Regression

✓ Divide the dataset into the training (70%) and the test (30%) at random

```
# Split the data into the training/test sets
trn_idx <- sample(1:1029, round(0.7*1029))
trn_data <- RegData[trn_idx,]
test_data <- RegData[-trn_idx,]
```

Data	
▶ concrete	1029 obs. of 9 variables
▶ RegData	1029 obs. of 9 variables
RegX	num [1:1029, 1:8] 2.485 0.495 0.495 -0.79 -0.143 ...
▶ test_data	309 obs. of 9 variables
▶ trn_data	720 obs. of 9 variables

# R Exercise

- k-NN Regression

✓ Find the best k based on LOO validation

```
# Find the best k using leave-one-out validation
nk <- c(1:10)
trn.n <- dim(trn_data)[1]
trn.v <- dim(trn_data)[2]

val.rmse <- matrix(0,length(nk),1)

for (i in 1:length(nk)){

  cat("k-NN regression with k:", nk[i], "\n")
  tmp_residual <- matrix(0,trn.n,1)

  for (j in 1:trn.n){

    # Data separation for leave-one-out validation
    tmptrnX <- trn_data[-j,1:(trn.v-1)]
    tmptrnY <- trn_data[-j,trn.v]
    tmpvalX <- trn_data[j,1:(trn.v-1)]
    tmpvalY <- trn_data[j,trn.v]

    # Train k-NN & evaluate
    tmp.knn.reg <- knn.reg(tmptrnX, test = tmpvalX, tmptrnY, k=nk[i])
    tmp_residual[j,1] <- tmpvalY - tmp.knn.reg$pred

  }

  val.rmse[i,1] <- sqrt(mean(tmp_residual^2))
}
```

```
> val.rmse
      [,1]
[1,] 9.151637
[2,] 8.897865
[3,] 8.524996
[4,] 8.838078
[5,] 9.019768
[6,] 9.087615
[7,] 9.212508
[8,] 9.458858
[9,] 9.441347
[10,] 9.507920
```

# R Exercise

- k-NN Regression

- ✓ Use the best k for the final model

```
# find the best k
best.k <- nk[which.min(val.rmse)]

# Evaluate the k-NN with the test data
test.knn.reg <- knn.reg(trn_data[,1:ncol(trn_data)-1], test = test_data[,1:ncol(test_data)-1],
                      trn_data[,ncol(trn_data)], k=best.k)

tgt.y <- test_data[,ncol(trn_data)]
knn.haty <- test.knn.reg$pred
```

- ✓ Train the MLR for benchmark

```
# Train the MLR for comparison
full_model <- lm(RegY ~ ., data = trn_data)
mlr.haty <- predict(full_model, newdata = test_data)
```

# R Exercise

- k-NN Regression

✓ Compare the performance of k-NN and MLR

```
# Regression performance comparison in terms of MAE
mean(abs(tgt.y-knn.haty))
mean(abs(tgt.y-mlr.haty))

# Plot the result
plot(tgt.y, knn.haty, pch = 1, col = 1, xlim = c(0,80), ylim = c(0, 80))
points(tgt.y, mlr.haty, pch = 2, col = 4, xlim = c(0,80), ylim = c(0,80))
abline(0,1,lty=3)
```

```
> mean(abs(tgt.y-knn.haty))
[1] 6.949288
> mean(abs(tgt.y-mlr.haty))
[1] 8.239489
```

# R Exercise

- k-NN Regression
  - ✓ Compare the performance of k-NN and MLR

